

FAUZAN EJAZ

+39 342 101 4345 | ejazfauzan14@gmail.com | Milan, Italy
github.com/fauzan111 | analyze.withfauzan.com | Italian B1 · English C1

PROFILE

Data Scientist with 4+ years of combined experience specialising in energy sector ML, time series forecasting, and production pipeline engineering. Direct production experience at GEKO S.p.A. delivering electricity consumption forecasting models and automated data pipelines for the Italian energy market. Proficient in Python, XGBoost, Scikit-learn, Optuna, MLflow, Docker, and CI/CD — with working knowledge of LSTM, Prophet, and anomaly detection for temporal energy data. Familiar with ARERA regulatory frameworks and Italian energy market data standards. Italian B1 and English C1 — a significant differentiator for roles at Italian utilities and European energy firms.

TECHNICAL SKILLS

ML & Forecasting	Scikit-learn · XGBoost · Optuna · LSTM / GRU · Prophet · ARIMA / SARIMA · Anomaly Detection · Feature Engineering · SHAP (XAI) · Imbalanced Classification
MLOps & Cloud	MLflow · Docker · CI/CD (GitHub Actions) · Weights & Biases · FastAPI · ONNX · AWS (CLF-C02 in progress) · Streamlit
Data Engineering	Python · SQL · Apache Spark · ETL Pipelines · Pandas · NumPy · Power BI · Time Series Processing
Deep Learning	PyTorch · Hugging Face Transformers · BERT · RoBERTa · YOLOv8 · LangGraph
EU / Energy	GDPR by Design · EU AI Act Compliance · ARERA Regulatory Frameworks · Smart Meter Data · Energy Consumption Modelling · Responsible AI

PROFESSIONAL EXPERIENCE

Data Scientist — Energy Forecasting & ML Systems

July 2025 – Present

GEKO S.p.A. (Energy & Ambient) · Naples, Italy

- Engineered a production-grade ETL pipeline processing [~X GWh / X million records] of Italian smart meter data daily, with multi-layer validation achieving 99% data integrity before model ingestion.
- Developed and deployed electricity consumption forecasting models (XGBoost, Scikit-learn) with Optuna hyperparameter optimisation, achieving [~X% MAPE / X% improvement] on seasonal Italian energy demand patterns.
- Built automated model evaluation and monitoring pipeline integrated into CI/CD (GitHub Actions + Docker), enabling real-time performance tracking and early drift detection across live forecasting outputs.
- Designed and executed unit and integration tests for Python-based forecasting modules, identifying and resolving 20+ edge-case bugs pre-deployment — ensuring model reliability in a regulated energy environment.
- Collaborated with domain stakeholders on ARERA-compliant data handling and energy market reporting, contributing to regulatory data governance workflows.

Data Analyst

Jan 2021 – Jan 2024

LTIMindtree · Kolkata, India

- Authored complex SQL pipelines for data extraction and transformation, reducing manual reporting preparation time by ~30% through automation.
- Built and maintained Power BI dashboards for KPI monitoring used daily by cross-functional teams, replacing manual spreadsheet workflows with live data connections.
- Conducted exploratory data analysis on large structured datasets, identifying data quality issues and correcting upstream ETL logic.

ML & DATA SCIENCE PROJECTS

Electricity Consumption Anomaly Detection (Energy Domain) XGBoost · LSTM · Prophet · MLflow · Python · SHAP

Time series anomaly detection system for Italian smart meter data, combining statistical baselines (ARIMA/Prophet) with gradient boosting (XGBoost) and LSTM deep learning for sequential pattern recognition. SHAP values used for model explainability — identifying which temporal features (hour-of-day, temperature, calendar effects) drove forecast deviations. MLflow for experiment tracking and model registry. Aligned with ARERA smart meter data standards. [Add GitHub link if public]

GDPR-Guardian — EU Regulatory Compliance Service [Microsoft Presidio](#) · [spaCy](#) · [FastAPI](#) · [Docker](#) · [EU AI Act](#) · [GDPR](#)

Production-ready PII detection and redaction API built for the Italian/European regulatory market. Custom Microsoft Presidio recognisers detect Codice Fiscale and Italian-specific sensitive entities using spaCy's Italian NLP model. FastAPI gateway with full Docker containerisation for scalable deployment. Directly aligned with GDPR by Design requirements and EU AI Act Article 10 data governance mandates. Demonstrates regulatory AI expertise relevant to Italian utility and public sector environments.

End-to-End Visual Quality Control System [YOLOv8](#) · [PyTorch](#) · [ONNX](#) · [MLflow](#) · [Streamlit](#)

Manufacturing visual inspection pipeline using YOLOv8 and PyTorch, achieving 78.9% mAP on MVTec AD dataset. ONNX export for CPU edge deployment with significantly reduced inference latency. Versioned data pipeline with automated mask-to-bounding-box conversion. Real-time monitoring via Streamlit and MLflow experiment tracking. Demonstrates full ML lifecycle ownership from data engineering to production deployment.

Sentiment Analysis — BERT/RobERTa Fine-tuning [Hugging Face Transformers](#) · [BERT](#) · [PyTorch](#) · [MLflow](#) · [Weights & Biases](#)

Fine-tuned BERT and RobERTa on 360K+ text samples using Hugging Face Transformers and PyTorch. Achieved 88% F1-score on 3-class sentiment classification, outperforming VADER and baseline models. Handled severe class imbalance and multilingual input. Full experiment tracking with MLflow and Weights & Biases. Docker-ready inference endpoint with real-time monitoring.

EDUCATION

MSc in Data Science 2023 – 2026
[University of Naples Federico II, Italy](#)
Fastweb + Vodafone 5G Academy Programme participant

Bachelor of Computer Applications 2018 – 2021
[Asansol Engineering College, India](#)
Grade: 9.0 / 10.0

CERTIFICATIONS

- **AWS Certified Cloud Practitioner (CLF-C02)** — Amazon Web Services *In progress — expected 2026*
- **EU AI Act Compliance Fundamentals** — EAIF / Self-study *In progress*

LANGUAGES

Italian B1 · **English** C1 · **Hindi** C1 · **Urdu** C1 · **Bengali** B1